

Topic: Human Activity Recognition using Recurrent Neural Networks

Domain: Sensor-Based Time-Series Classification

Dataset: UCI Human Activity Recognition Using Smartphones

Course: Machine Learning

Assignment: 5.4

Group Members:

- Eman Shahbaz (BITF22M516)
- Dua Sarwar (BITF22M530)

1. Objective

The objective of this assignment is to implement **Recurrent Neural Network (RNN)-based classification** for time-series sensor data. Unlike previous deliverables that relied on engineered or flattened features, this assignment focuses on preserving **temporal dependencies** in raw sensor signals. Three recurrent architectures were implemented and evaluated:

- Simple RNN
- Long Short-Term Memory (LSTM)
- Gated Recurrent Unit (GRU)

2. Dataset Description

The **UCI HAR Dataset** contains recordings from a smartphone's accelerometer and gyroscope sensors. The phone was worn on the waist while subjects performed daily activities.

Sensors used:

- 3-axis accelerometer (x, y, z)
- 3-axis gyroscope (x, y, z)

Activities (6 classes):

1. Walking
2. Walking Upstairs
3. Walking Downstairs

4. Sitting
5. Standing
6. Lying

Dataset size:

- Training samples: 7,352
- Testing samples: 2,947

Each sample is a **time window of 128 timesteps**.

3. Preprocessing for Sequential Data

For recurrent models, **raw inertial signals** were used from the *Inertial Signals* directory instead of pre-extracted feature vectors.

Each sample was represented as a **3D tensor**:

$(\text{samples}, \text{timesteps}, \text{features}) = (N, 128, 6)$

Where:

- Timesteps = 128
- Features = 6 sensor channels (acc_x, acc_y, acc_z, gyro_x, gyro_y, gyro_z)

Normalization

Signals were normalized **per channel** using training-set mean and standard deviation to ensure stable learning and avoid scale imbalance.

4. Recurrent Models Implemented

Three recurrent architectures were implemented using TensorFlow/Keras:

4.1 Simple RNN

- Single recurrent layer
- Suffers from vanishing gradient problem
- Serves as a baseline recurrent model

4.2 LSTM

- Uses input, forget, and output gates
- Designed to capture long-term temporal dependencies
- More stable than Simple RNN

4.3 GRU

- Simplified gated architecture
- Fewer parameters than LSTM
- Faster convergence and good generalization

All models used:

- Adam optimizer
- Sparse categorical cross-entropy loss
- Softmax output layer for 6-class classification

5. Evaluation Metrics

For **each model**, the following metrics were reported:

- Accuracy
- Weighted F1-score
- Confusion Matrix
- Training vs Validation Loss Curves

6. Results

6.1 RNN Model Comparison

Model	Accuracy	Weighted F1
Simple RNN	0.4404	0.3866
LSTM	0.6339	0.5590
GRU	0.6624	0.6179

Best Recurrent Model: GRU

6.2 Training vs Validation Behavior

- **Simple RNN** showed unstable convergence and underfitting, with poor validation performance.
- **LSTM** demonstrated smoother training and improved validation loss.
- **GRU** achieved the most stable learning curves, indicating better generalization and efficient learning.

7. Comparison with Fully Connected Neural Network

Model	Accuracy
FCNN (Assignment 5.3)	0.9437
Best RNN (GRU)	0.6624

Accuracy Difference:

$$0.6624 - 0.9437 = -0.2813$$

Interpretation

The FCNN outperformed all recurrent models because it was trained on **engineered, window-based features** that already encode discriminative motion patterns. In contrast, RNN-based models operate on **raw sensor signals**, which are noisier and require more complex temporal modeling and tuning.

8. Limitations of Recurrent Models

- Simple RNN suffers from vanishing gradients.
- LSTM and GRU increase computational cost.
- Training time is longer compared to FCNN.
- Performance is sensitive to hyperparameter tuning.
- RNNs may underperform when datasets already provide strong engineered features.

9. Conclusion

This assignment successfully implemented and evaluated **Simple RNN**, **LSTM**, and **GRU** models for time-series classification using raw sensor data. GRU achieved the best recurrent performance, while FCNN from Assignment 5.3 remained superior due to richer feature representations. The results highlight the trade-offs between sequential modeling and feature-based learning.